






V-JEPA predictor's L1 on masked next-frame tokens — the JEPA objective.
Should TREND DOWN. If it climbs ⇒ encoder forgetting predictive structure.