


future_mse — sorted ↑ lowest first · 8 encoders · 95 % BCa CI

future_mse — sorted ↑ lowest first



V-JEPA predictor L1 on masked next-frame tokens (the JEPA objective; V-JEPA-only).

Layman: e.g. given the first frames, how well does the model's predictor guess the next frame's content?

TEST split · whiskers = BCa 95 % CI from 10 K bootstrap.