


taxonomy_f1 — sorted ↓ highest first · 24 encoders · 95 % BCa CI

taxonomy_f1 — sorted ↓ highest first



Mean over 15 VLM-tag dims of per-dim top-1 (single) / sample-F1 (multi).
Layman: e.g. can a linear read-off name scene attributes — crowd density, lighting, camera motion — from the features?
TEST split · whiskers = BCa 95 % CI from 10 K bootstrap.