


teacher\_free — sorted ↑ lowest first · 16 encoders · 95 % BCa CI

teacher\_free — sorted ↑ lowest first



Free-running minus teacher-forced L1 — lower = less exposure bias.  
Layman: e.g. how much worse is the model when it must rely on its OWN predictions vs being fed the truth each step?  
TEST split · whiskers = BCa 95 % CI from 10 K bootstrap.