


0 · Experiment map — 13 arms, one backbone (ViT-G 2B)




0 · Techniques — FULL FORM + like you're 5

technique	FULL FORM (glossary)	like you're 5
frozen	frozen backbone (no adaptation)	Don't touch the brain; just test what it knows.
pretrain (vCSSL)	vanilla continual SSL (Self-Supervised Learning) — OURS anchor	Keep practicing videos the same old way.
surgery 3stage-DI (OURS)	staged factor-curriculum continual-FT, with D_I tubes	Special lesson clips; unlock slowly; includes things-touching clips.
surgery noDI (OURS)	staged factor-curriculum continual-FT, without D_I	Same slow unlocking; skip the things-touching clips.
surgery heads (OURS)	staged factor-curriculum continual-FT, head-only	Brain stays locked; only a tiny helper learns.
surgery RAW (control)	staged factor-curriculum continual-FT on raw clips	Surgery's slow unlocking, but with plain videos.
Auto-RGN	Automatic Relative Gradient Norm (Lee et al. ICLR'23)	Layers that pull harder get bigger learning steps.
Full-FT	Full Fine-Tuning	Unlock the whole brain; change everything at once.
LP-FT	Linear-Probing then Fine-Tuning (Kumar et al. ICLR'22)	First teach the helper, then unlock everything.
LoRA	Low-Rank Adaptation (Hu et al. 2021)	Stick tiny add-on notes; never rewrite the book.
DoRA	Weight-Decomposed Low-Rank Adaptation (Liu et al. 2024)	Sticky notes plus a volume knob per page.
CaSSLe	continual self-supervised distillation (Fini et al. CVPR'22)	New learning must still match the old teacher.
EWC	Elastic Weight Consolidation (Kirkpatrick et al. PNAS'17)	Important old memories get glued — hard to change.

1 · HERO — kept-checkpoint scorecard

iter18 KEPT-checkpoint scorecard — the encoder each arm EXPORTS to eval




VERDICT — best OURS (surgery; black-edged bars) vs best OTHER · whiskers = 95% CI (per-clip BCa; legacy top-1 rows: binomial approx; causal/mask-ratio/val-loss have no per-clip data → no whisker)

action top-1 OURS 0.2927 · other 0.3348 = TIE (CIs overlap)	motion-cos OURS 0.1573 · other 0.2545 OTHERS LEAD	future L1 OURS 0.4995 · other 0.5011 OURS LEAD	causal L1 OURS 0.5296 · other 0.5301 OURS LEAD	mask-ratio OURS 0.0547 · other 0.0528 OTHERS LEAD	val JEPA loss OURS 0.1968 · other 0.4461 OURS LEAD
--	--	---	---	--	---

- Surgery leads future-L1 (0.4995) and causal (0.530); LoRA/DoRA trail ≤ 0.005 ; full-FT owns semantics.
- Semantic gaps sit inside N=451 noise; surgery's ~ 0.035 future-L1 lead over pretrain/full-FT holds.

2 · Training — every probe checkpoint, every arm

iter18 TRAIN probe checkpoints — every arm, every probe · star marker = KEPT ckpt · OURS = greens (solid enc / dotted head)



- Surgery's future-L1 falls every stage (0.523→0.4995); pretrain stays flat (~0.54) — progressive unfreeze works.
- Selector keeps mid-training minima — LP-FT probe 1/5, full-FT 2/4; last checkpoints regress.

3 · Upcoming — per-encoder TEST eval scorecard

**no EVAL artifacts yet —
per-encoder evals launch as each arm finishes training**

- TEST verdicts (n=1,825, 95% BCa CIs) auto-fill as evals finish; val trends need confirmation.
- Paired-delta finale decides significance — val-probe leads are directional only.


4a · iter17 — frozen backbone selection

FROZEN encoders — ABSOLUTE head-metric values (no Δ) · value $\pm$ BCa 95% CI · colour = per-metric min-max (green = best) · BLUE BOX = best per metric
10 frozen encoders $\times$ 3 shared metrics · sorted by action_top1 · 10 K bootstrap

	action_top1 ↑ better	motion_cos ↑ better	taxonomy_f1 ↑ better
V-JEPA 2.1 frozen	44.384 ± 2.274	0.009 ± 0.001	0.793 ± 0.016
vjepa 2 1 vitL frozen	44.219 ± 2.301	0.004 ± 0.000	0.788 ± 0.017
vjepa 1 vitH frozen	40.493 ± 2.219	0.007 ± 0.001	0.702 ± 0.019
lejepa vitL frozen	40.055 ± 2.219	0.014 ± 0.002	0.740 ± 0.018
vjepa 1 vitL frozen	39.890 ± 2.274	0.008 ± 0.001	0.660 ± 0.019
ijepa vitH14	39.123 ± 2.247	0.016 ± 0.001	0.781 ± 0.017
vjepa 2 0 vitg ssv2	38.849 ± 2.274	0.007 ± 0.001	0.776 ± 0.017
DINOv2 frozen	38.466 ± 2.247	0.016 ± 0.001	0.816 ± 0.016
vjepa 2 vitL 256 frozen	37.918 ± 2.247	0.013 ± 0.001	0.778 ± 0.017
ijepa vitG16	37.480 ± 2.219	0.019 ± 0.001	0.787 ± 0.017

- V-JEPA 2.1 ViT-G is the strongest frozen backbone — justified as iter18's sole backbone.
- Every frozen model is motion-blind (motion-cos ~ 0) — adaptation is necessary, not optional.

4b · iter17 — ViT-G family raw values after training



- Any continual training beats frozen by +6-7pp top-1; gains are not surgery-specific.
- Only surgery moves future-MSE and causal — encoder updates drive temporal gains, heads don't.

4c · iter17 — paired deltas with 95% CI

PAIRED surgery – vanilla cont-SSL (champion duel) across backbones · per cell: Δ + 95% BCa CI
 deep green = surgery sig. better · deep red = vanilla cont-SSL · pale grey = TIE (95% CI crosses 0) · 10 K bootstrap

action_top1	-1.37 [-1.12, +0.329] TIE	-0.438 [-2.30, +1.48] TIE	+0.219 [-1.26, +1.81] TIE
motion_cos	+0.010 [+0.0083, +0.012] SURGERY	+0.018 [+0.017, +0.020] SURGERY	-0.102 [-0.103, -0.099] VANILLA CONT-SSL
taxonomy_f1	-0.0001 [-0.015, +0.015] TIE	-0.0085 [-0.024, +0.0066] TIE	-0.0080 [-0.025, +0.0089] TIE
future_mse	+0.025 [+0.024, +0.025] SURGERY	+0.0037 [+0.0030, +0.0044] SURGERY	-0.070 [-0.072, -0.067] VANILLA CONT-SSL
rollout	0 = equal TIE	-0.0001 [-0.0002, -0.0000] VANILLA CONT-SSL	0 = equal TIE
causal	+0.019 [+0.018, +0.020] SURGERY	+0.0012 [+0.0008, +0.0013] SURGERY	-0.062 [-0.063, -0.060] VANILLA CONT-SSL
tdist	0 = equal TIE	-0.0009 [-0.0011, -0.0008] VANILLA CONT-SSL	0 = equal TIE
teacher_free	0 = equal TIE	+0.0014 [+0.0011, +0.0017] SURGERY	-0.0019 [-0.0022, -0.0016] VANILLA CONT-SSL
maskratio	+0.0091 [+0.0088, +0.0094] SURGERY	+0.0018 [+0.0016, +0.0021] SURGERY	-0.0078 [-0.0083, -0.0070] VANILLA CONT-SSL
WINS	Surgery 4 Vanilla cont-SSL 0 tie 5 V-JEPA 2.1 VIT-G ~2B params	Surgery 5 Vanilla cont-SSL 2 tie 2 V-JEPA 2.1 VIT-g ~1B params	Surgery 0 Vanilla cont-SSL 5 tie 4 V-JEPA 2.0 VIT-g ~1B params

- Surgery beats frozen significantly on action and motion-cos — CIs exclude zero.
- Versus pretrain: action ties; motion/prediction edge persists — surgery's value is temporal, not semantic.