




V-JEPA predictor L1 on masked next-frame tokens (the JEPA objective; V-JEPA-only).
Layman: e.g. given the first frames, how well does the model's predictor guess the next frame's content?
TEST split · whiskers = BCa 95 % CI from 10 K bootstrap.