


iter18 KEPT-checkpoint scorecard — the encoder each arm EXPORTS to eval



VERDICT — best OURS (surgery; black-edged bars) vs best OTHER · whiskers = 95% CI (per-clip BCa; legacy top-1 rows: binomial approx; causal/mask-ratio/val-loss have no per-clip data → no whisker)

action top-1
OURS 0.2927 · other 0.3348
≈ TIE (CIs overlap)

motion-cos
OURS 0.1573 · other 0.2545
OTHERS LEAD

future L1
OURS 0.4995 · other 0.5011
OURS LEAD

causal L1
OURS 0.5296 · other 0.5302
OURS LEAD

mask-ratio
OURS 0.0547 · other 0.0568
OURS LEAD

val JEPA loss
OURS 0.1968 · other 0.4463
OURS LEAD