


iter18 KEPT-checkpoint scorecard — the encoder each arm EXPORTS to eval



VERDICT — best OURS (surgery; black-edged bars) vs best OTHER · whiskers = 95% CI (per-clip BCa; legacy top-1 rows: binomial approx; causal/mask-ratio/val-loss have no per-clip data → no whisker)

action top-1 OURS 0.2927 · other 0.3348 ≈ TIE (CIs overlap)	motion-cos OURS 0.1501 · other 0.2545 OTHERS LEAD	future L1 OURS 0.4995 · other 0.5011 OURS LEAD	causal L1 OURS 0.5296 · other 0.5301 OURS LEAD	mask-ratio OURS 0.0547 · other 0.0528 OTHERS LEAD	val JEPA loss OURS 0.4465 · other 0.4461 OTHERS LEAD
--	--	---	---	--	---