



iter18 KEPT-checkpoint scorecard — the encoder each arm EXPORTS to eval


action top-1 ↑ · best: fullFT




motion-cos ↑ · best: fullFT




future L1 ↓ · best: s3DI-enc




causal L1 ↓ · best: sNoDI-enc



mask-ratio ↓ · best: DoRA



val JEPa loss ↓ · best: s3DI-hd (hd)



VERDICT — best OURS (surgery; black-edged bars) vs best OTHER · whiskers = 95% CI (per-clip BCa; legacy top-1 rows: binomial approx; causal/mask-ratio/val-loss have no per-clip data → no whisker)

action top-1
OURS 0.2927 · other 0.3348
≈ TIE (CIs overlap)

motion-cos
OURS 0.1573 · other 0.2545
OTHERS LEAD

future L1
OURS 0.4995 · other 0.5011
OURS LEAD

causal L1
OURS 0.5296 · other 0.5301
OURS LEAD

mask-ratio
OURS 0.0547 · other 0.0528
OTHERS LEAD

val JEPa loss
OURS 0.1968 · other 0.4461
OURS LEAD