

V-JEPA 2.1
ViT-G (2B)
= BIGGEST avail.
checkpoints/
vjepa2_1_vitG_384.pt

adaptation family

PROPOSED

0.3-frozen

B1 PEFT
(Parameter-Efficient
Fine-Tuning)

B2 Surgical-FT

B3 Cont-SSL

B4(a) Full-FT
(Full Fine-Tuning) /
B4(b) LP-FT
(Linear-Probing
then Fine-Tuning)

CONTROL

0.7·S

m09c1 surgery
factor curriculum
D_L→D_A→D_I
SALT + SPD +
saliency + replay

★ iter18 improvements
(variants of S ·
plan_outperform_FT.md)

★ wiseft
src/utls/wiseft_merge.py
post-hoc
0.7·S + 0.3·frozen

m09b_peft.py
LoRA
(Low-Rank Adaptation)
→ DoRA
(Weight-Decomposed
Low-Rank Adaptation)
adapters on
attn.qkv + mlp (r=16)

m09c1 + auto_rgn
freeze
RGN = ||g_block|| /
||θ_block||
top-k blocks,
RAW clips

m09a1 + CaSSLe
distill
(FROZEN teacher) +
EWC (Elastic Weight
Consolidation)
Fisher anchor,
RAW clips

m09a1
B4(a):
unfreeze_below=1.0
(full)
B4(b): lp-ft-stage0
then unfreeze

m09c1 surgery
factors OFF · RAW
= surgery_raw
(the ablation)

📄 replay25
cfg: raw replay
0.5 → 0.25

⬆️ diheavy
cfg: D_I stage
0.30 → 0.45

+ tcclaux
cfg: gamma_tcc 0 → 0.1
reuse γ·TCC loss
(Dwibedi CVPR'19)

+ intervene
cfg: +object-tube mask
(Causal-JEPA
arXiv:2602.11389)

student_encoder.pt
(+ predictor)
export_student_
for_eval()

configs/eval/
probe_encoders.yaml
(one row per
baseline×arm)

run_eval.sh →
m12a..f
14 metrics +
BCa 95% CI

m13 \$G hero table
surgery vs 4
baselines vs anchors