


iter18 KEPT-checkpoint scorecard — the encoder each arm EXPORTS to eval



VERDICT — best OURS (surgery; black-edged bars) vs best OTHER · whiskers = 95% CI (per-clip BCa; legacy top-1 rows: binomial approx; causal/mask-ratio/val-loss have no per-clip data $\rightarrow$ no whisker)

