


teacher_free — sorted ↑ lowest first · 14 encoders · 95 % BCa CI



Free-running minus teacher-forced L1 — lower = less exposure bias.

Layman: e.g. how much worse is the model when it must rely on its OWN predictions vs being fed the truth each step?

TEST split · whiskers = BCa 95 % CI from 10 K bootstrap.