



probe_taxonomy — per-dim metric across 14 encoders · 95 % BCa CI
(top-1 for single-label, sample-F1 for multi-label)


crowd_density
(top1_acc, n=3cls)




infrastructure_quality
(top1_acc, n=3cls)




lighting
(top1_acc, n=4cls)




notable_objects
(sample_f1, n=15cls)




pedestrian_vehicle_separation
(top1_acc, n=3cls)




road_encroachment
(top1_acc, n=3cls)




road_layout
(sample_f1, n=8cls)




road_surface
(top1_acc, n=8cls)




scene_type
(top1_acc, n=13cls)




time_of_day
(top1_acc, n=2cls)




traffic_density
(top1_acc, n=3cls)




traffic_mix
(top1_acc, n=4cls)



vegetation
(top1_acc, n=4cls)



video_quality
(top1_acc, n=3cls)



weather
(top1_acc, n=5cls)

