




m09c · $\lambda=0.005$ · Probe trajectory — encoder evolution across SSL training
(VAL split · N=451 probe-clips · per-step diagnostic, NOT paper-final eval)



AttentiveClassifier accuracy on K-class motion-flow probe at each val cycle.
Should TREND UP as the encoder learns to separate motion classes.



Same-class minus different-class cosine separation.
> 0 ⇒ semantic clustering; near 0 ⇒ encoder doesn't separate motion classes.



V-JEPA predictor's L1 on masked next-frame tokens — the JEPA objective.
Should TREND DOWN. If it climbs ⇒ encoder forgetting predictive structure.