


iter18 KEPT-checkpoint scorecard — the encoder each arm EXPORTS to eval



VERDICT — best OURS (surgery; black-edged bars) vs best OTHER · whiskers = 95% CI (per-clip BCa; legacy top-1 rows: binomial approx; causal/mask-ratio/val-loss have no per-clip data → no whisker)

action top-1 (awaiting data)	motion-cos (awaiting data)	future L1 (awaiting data)	causal L1 (awaiting data)	mask-ratio (awaiting data)	val JEPa loss (awaiting data)
---------------------------------	-------------------------------	------------------------------	------------------------------	-------------------------------	----------------------------------