


FactorJEPA vs. the best competitor, across scale: separation in 95%-CI units
 (rows sorted best-to-worst by the FULL 1B panel; dashed line = the 1× separation threshold)



green: separated win ($\geq 1\times$) yellow: ahead but within CI (0-1×) red: competitor ahead (<0)

$n_{\text{total}} = \text{train} + \text{val} + \text{test}$. FULL ViT-G 2B omitted: not run under the compute budget ($\approx 2\times$ the 1B cost).